

Stock Price Prediction Project

```
from platform import python_version
print(python_version())
```

3.9.12

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Importing Data

```
stock_df=pd.read_csv('Stock_Price_data_set.csv')
```

stock_df

	Date	Open	High	Low	Close	Adj
0	2018-02-05	262.000000	267.899994	250.029999	254.259995	
1	2018-02-06	247.699997	266.700012	245.000000	265.720001	
2	2018-02-07	266.579987	272.450012	264.329987	264.559998	
3	2018-02-08	267.079987	267.619995	250.000000	250.100006	
4	2018-02-09	253.850006	255.800003	236.110001	249.470001	
...	
1004	2022-01-31	401.970001	427.700012	398.200012	427.140015	
1005	2022-02-01	432.959991	458.480011	425.540009	457.130005	
1006	2022-02-02	448.250000	451.980011	426.480011	429.480011	
1007	2022-02-03	421.440002	429.260010	404.279999	405.600006	
1008	2022-02-04	407.309998	412.769989	396.640015	410.170013	
...	
Volume						
0	11896100					
1	12595800					
2	8981500					
3	9306700					
4	16906900					
...	...					
1004	20047500					
1005	22542300					
1006	14346000					

```
1007    9905200
1008    7782400
```

```
[1009 rows x 7 columns]
```

Exploring Data

```
stock_df.shape
```

```
(1009, 7)
```

```
stock_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1009 entries, 0 to 1008
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	Date	1009 non-null	object
1	Open	1009 non-null	float64
2	High	1009 non-null	float64
3	Low	1009 non-null	float64
4	Close	1009 non-null	float64
5	Adj Close	1009 non-null	float64
6	Volume	1009 non-null	int64

```
dtypes: float64(5), int64(1), object(1)
```

```
memory usage: 55.3+ KB
```

```
stock_df.columns
```

```
Index(['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume'],  
      dtype='object')
```

```
stock_df.describe
```

```
<bound method NDFrame.describe of
```

	Date	Open	High	Low	Close	Adj Close
0	2018-02-05	254.259995	262.000000	247.899994	250.029999	254.259995
1	2018-02-06	265.720001	247.699997	266.700012	245.000000	265.720001
2	2018-02-07	264.559998	266.579987	272.450012	264.329987	264.559998
3	2018-02-08	250.100006	267.079987	267.619995	250.000000	250.100006
4	2018-02-09	249.470001	253.850006	255.800003	236.110001	249.470001
...
1004	2022-01-31	427.140015	401.970001	427.700012	398.200012	427.140015
1005	2022-02-01	457.130005	432.959991	458.480011	425.540009	457.130005

```

457.130005
1006  2022-02-02  448.250000  451.980011  426.480011  429.480011
429.480011
1007  2022-02-03  421.440002  429.260010  404.279999  405.600006
405.600006
1008  2022-02-04  407.309998  412.769989  396.640015  410.170013
410.170013

```

```

      Volume
0      11896100
1      12595800
2       8981500
3       9306700
4      16906900
...      ...
1004    20047500
1005    22542300
1006    14346000
1007     9905200
1008     7782400

```

```
[1009 rows x 7 columns]>
```

Missing Values

```
stock_df.isna().any()
```

```

Date          False
Open          False
High          False
Low           False
Close         False
Adj Close     False
Volume        False
dtype: bool

```

Duplicates

```
stock_df.duplicated().sum()
```

```
0
```

Column Data Type

```
stock_df.dtypes
```

```

Date          object
Open          float64
High          float64
Low           float64
Close         float64
Adj Close     float64

```

```
Volume          int64
dtype: object
```

Outliers

```
plt.subplot(2,3,1)
stock_df['Open'].plot(kind='box')
```

```
plt.subplot(2,3,2)
stock_df['Close'].plot(kind='box')
```

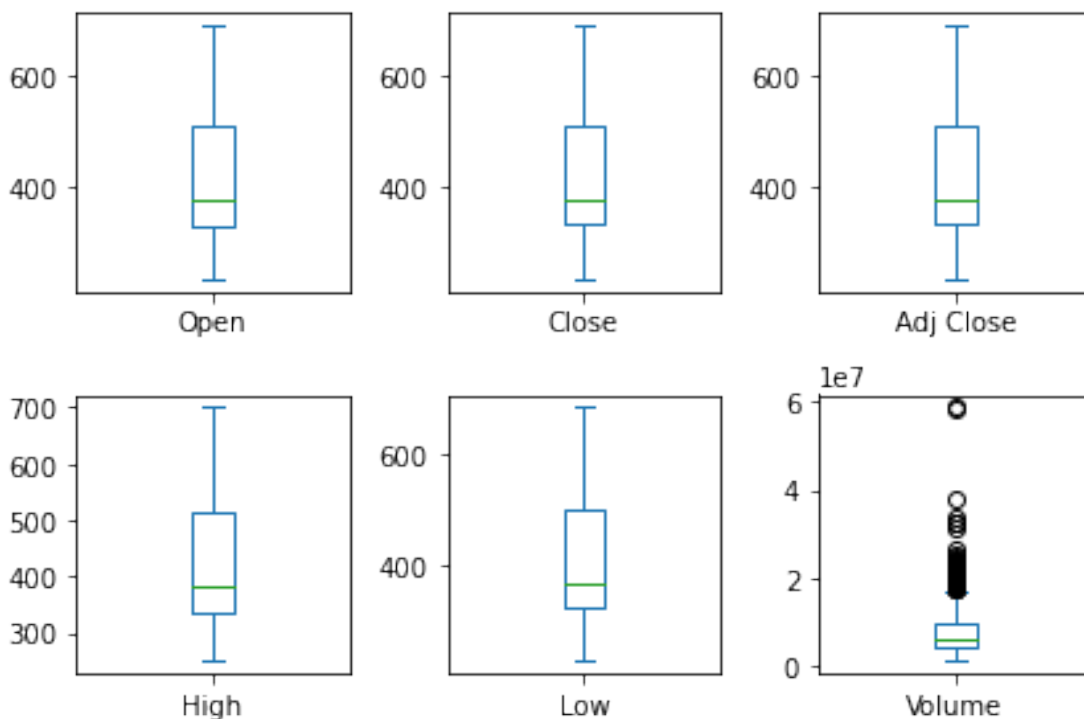
```
plt.subplot(2,3,3)
stock_df['Adj Close'].plot(kind='box')
```

```
plt.subplot(2,3,4)
stock_df['High'].plot(kind='box')
```

```
plt.subplot(2,3,5)
stock_df['Low'].plot(kind='box')
```

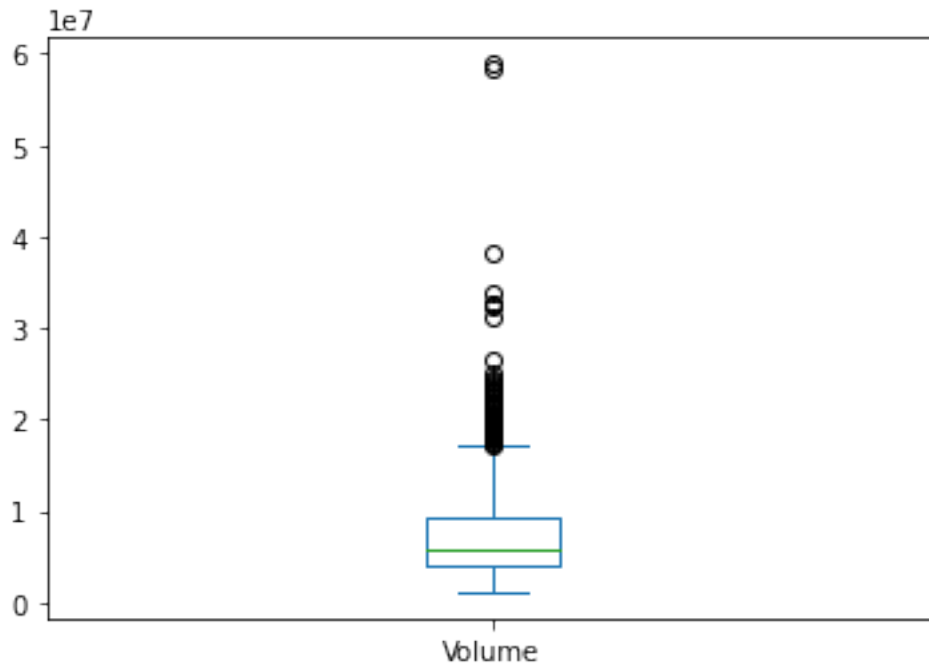
```
plt.subplot(2,3,6)
stock_df['Volume'].plot(kind='box')
```

```
plt.tight_layout()
```



```
stock_df['Volume'].plot(kind='box')
```

<AxesSubplot:>



```
def find_outlier_limits(col_name):
    Q1,Q3=stock_df[col_name].quantile([.25,.75])
    IQR=Q3-Q1
    low=Q1-(2* IQR)
    high=Q3+(2* IQR)
    return (high,low)

high_vol,low_vol=find_outlier_limits('Volume')
print('Volume: ','upper limit: ',high_vol,' lower limit: ',low_vol)

Volume:  upper limit:  19783400.0  lower limit:  -6369100.0

low_limit = 0
print('Volume: ','upper limit: ',high_vol,'lower limit: ',low_limit)

Volume:  upper limit:  19783400.0 lower limit:  0

#replacing outliers value
stock_df.loc[stock_df['Volume'] > high_vol,'Volume'] = high_vol

stock_df.loc[stock_df['Volume']>high_vol,'Volume']=high_vol

plt.subplot(2,3,1)
stock_df['Open'].plot(kind='box')

plt.subplot(2,3,2)
stock_df['Close'].plot(kind='box')

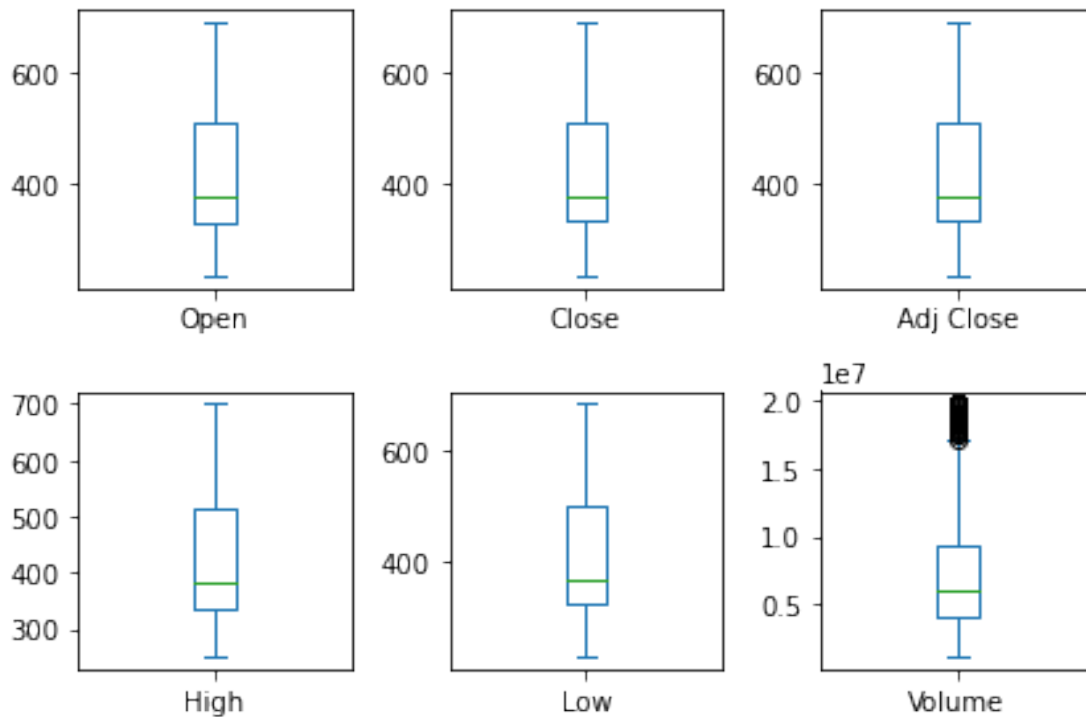
plt.subplot(2,3,3)
stock_df['Adj Close'].plot(kind='box')
```

```
plt.subplot(2,3,4)
stock_df['High'].plot(kind='box')


plt.subplot(2,3,5)
stock_df['Low'].plot(kind='box')

plt.subplot(2,3,6)
stock_df['Volume'].plot(kind='box')

plt.tight_layout()
```



```
stock_df['Volume'].plot(kind='box')
<AxesSubplot:>
```



```
outliers = [stock_df['Volume'] > high_vol, 'Volume']
outliers[True]
```

'Volume'

ML MODELING

stock_df

	Date	Open	High	Low	Close	Adj
0	2018-02-05	262.000000	267.899994	250.029999	254.259995	
1	2018-02-06	247.699997	266.700012	245.000000	265.720001	
2	2018-02-07	266.579987	272.450012	264.329987	264.559998	
3	2018-02-08	267.079987	267.619995	250.000000	250.100006	
4	2018-02-09	253.850006	255.800003	236.110001	249.470001	
...	
1004	2022-01-31	401.970001	427.700012	398.200012	427.140015	
1005	2022-02-01	432.959991	458.480011	425.540009	457.130005	
1006	2022-02-02	448.250000	451.980011	426.480011	429.480011	
1007	2022-02-03	421.440002	429.260010	404.279999	405.600006	

```
1008 2022-02-04 407.309998 412.769989 396.640015 410.170013
410.170013
```

```

      Volume
0      11896100
1      12595800
2       8981500
3       9306700
4      16906900
...
1004  19783400
1005  19783400
1006  14346000
1007   9905200
1008   7782400
```

```
[1009 rows x 7 columns]
```

```
X = stock_df.iloc[:, 1:8]
y = stock_df.iloc[:, 0]
print(X)
print(y)
```

	Open	High	Low	Close	Adj Close
Volume					
0	262.000000	267.899994	250.029999	254.259995	254.259995
11896100					
1	247.699997	266.700012	245.000000	265.720001	265.720001
12595800					
2	266.579987	272.450012	264.329987	264.559998	264.559998
8981500					
3	267.079987	267.619995	250.000000	250.100006	250.100006
9306700					
4	253.850006	255.800003	236.110001	249.470001	249.470001
16906900					
...
...					
1004	401.970001	427.700012	398.200012	427.140015	427.140015
19783400					
1005	432.959991	458.480011	425.540009	457.130005	457.130005
19783400					
1006	448.250000	451.980011	426.480011	429.480011	429.480011
14346000					
1007	421.440002	429.260010	404.279999	405.600006	405.600006
9905200					
1008	407.309998	412.769989	396.640015	410.170013	410.170013
7782400					

```
[1009 rows x 6 columns]
0      2018-02-05
```



```
1      2018-02-06
2      2018-02-07
3      2018-02-08
4      2018-02-09
```

```
...
1004    2022-01-31
1005    2022-02-01
1006    2022-02-02
1007    2022-02-03
1008    2022-02-04
```

```
Name: Date, Length: 1009, dtype: object
```

```
X = pd.get_dummies(X)
X
```

	Open	High	Low	Close	Adj Close
Volume					
0	262.000000	267.899994	250.029999	254.259995	254.259995
11896100					
1	247.699997	266.700012	245.000000	265.720001	265.720001
12595800					
2	266.579987	272.450012	264.329987	264.559998	264.559998
8981500					
3	267.079987	267.619995	250.000000	250.100006	250.100006
9306700					
4	253.850006	255.800003	236.110001	249.470001	249.470001
16906900					
...
...					
1004	401.970001	427.700012	398.200012	427.140015	427.140015
19783400					
1005	432.959991	458.480011	425.540009	457.130005	457.130005
19783400					
1006	448.250000	451.980011	426.480011	429.480011	429.480011
14346000					
1007	421.440002	429.260010	404.279999	405.600006	405.600006
9905200					
1008	407.309998	412.769989	396.640015	410.170013	410.170013
7782400					

```
[1009 rows x 6 columns]
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.2, random_state = 0)
```

```
from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
scale.fit_transform(X_train)
scale.transform(X_test);
```

```
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors=5,metric="euclidean")
```

TESTING

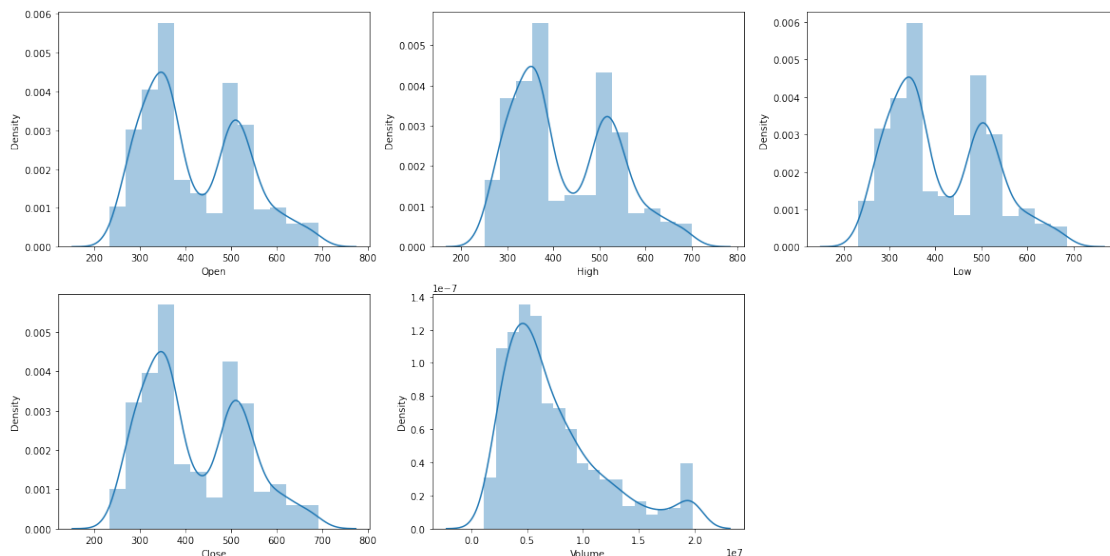
```
import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
#from xgboost import XGBClassifier
from sklearn import metrics
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
features = ['Open', 'High', 'Low', 'Close', 'Volume']
```

```
plt.subplots(figsize=(20,10))
```

```
for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sb.distplot(stock_df[col])
plt.show()
```



```
splitted = stock_df['Date'].str.split('-', expand=True)
```

```
stock_df['day'] = splitted[1].astype('int')
stock_df['month'] = splitted[0].astype('int')
stock_df['year'] = splitted[2].astype('int')
```

```
stock_df.head()
```

```

      Date      Open      High      Low      Close  Adj
Close \
```

```

0 2018-02-05 262.000000 267.899994 250.029999 254.259995
254.259995
1 2018-02-06 247.699997 266.700012 245.000000 265.720001
265.720001
2 2018-02-07 266.579987 272.450012 264.329987 264.559998
264.559998
3 2018-02-08 267.079987 267.619995 250.000000 250.100006
250.100006
4 2018-02-09 253.850006 255.800003 236.110001 249.470001
249.470001

```

```

      Volume  day  month  year
0  11896100    2   2018     5
1  12595800    2   2018     6
2   8981500    2   2018     7
3   9306700    2   2018     8
4  16906900    2   2018     9

```

```

stock_df['is_quarter_end'] = np.where(stock_df['month']%3==0,1,0)
stock_df.head()

```

```

      Date      Open      High      Low      Close  Adj
Close \
0 2018-02-05 262.000000 267.899994 250.029999 254.259995
254.259995
1 2018-02-06 247.699997 266.700012 245.000000 265.720001
265.720001
2 2018-02-07 266.579987 272.450012 264.329987 264.559998
264.559998
3 2018-02-08 267.079987 267.619995 250.000000 250.100006
250.100006
4 2018-02-09 253.850006 255.800003 236.110001 249.470001
249.470001

```

```

      Volume  day  month  year  is_quarter_end
0  11896100    2   2018     5                0
1  12595800    2   2018     6                0
2   8981500    2   2018     7                0
3   9306700    2   2018     8                0
4  16906900    2   2018     9                0

```

```

df=stock_df

```

```


data_grouped = df.groupby('year').mean()
plt.subplots(figsize=(20,10))

```

```

for i, col in enumerate(['Open', 'High', 'Low', 'Close']):
    plt.subplot(2,2,i+1)
    data_grouped[col].plot.bar()
plt.show()

```




```
df.groupby('is_quarter_end').mean()
```

	Open	High	Low	Close	Adj
Close \ is_quarter_end					
0	448.108608	454.981568	440.919864	448.071555	
1	341.911595	346.547464	336.562137	341.794528	

	Volume	day	month	year
is_quarter_end				
0	7.155014e+06	6.716235	2019.721692	15.785812
1	7.866719e+06	6.072464	2019.260870	15.536232

```
df['open-close'] = df['Open'] - df['Close']
df['low-high'] = df['Low'] - df['High']
df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0)
```

```
plt.pie(df['target'].value_counts().values,
        labels=[0, 1], autopct='%1.1f%%')
plt.show()
```



```
plt.figure(figsize=(10, 10))  
sb.heatmap(df.corr() > 0.9, annot=True, cbar=False)  
plt.show()
```

Open	1	1	1	1	1	0	0	0	0	0	0	0	0
High	1	1	1	1	1	0	0	0	0	0	0	0	0
Low	1	1	1	1	1	0	0	0	0	0	0	0	0
Close	1	1	1	1	1	0	0	0	0	0	0	0	0
Adj Close	1	1	1	1	1	0	0	0	0	0	0	0	0
Volume	0	0	0	0	0	1	0	0	0	0	0	0	0
day	0	0	0	0	0	0	1	0	0	0	0	0	0
month	0	0	0	0	0	0	0	1	0	0	0	0	0
year	0	0	0	0	0	0	0	0	1	0	0	0	0
is_quarter_end	0	0	0	0	0	0	0	0	0	1	0	0	0
open-close	0	0	0	0	0	0	0	0	0	0	1	0	0
low-high	0	0	0	0	0	0	0	0	0	0	0	1	0
target	0	0	0	0	0	0	0	0	0	0	0	0	1
	Open	High	Low	Close	Adj Close	Volume	day	month	year	is_quarter_end	open-close	low-high	target

```
features = df[['open-close', 'low-high', 'is_quarter_end']]
target = df['target']
```

```
scaler = StandardScaler()
features = scaler.fit_transform(features)
```

```
X_train, X_valid, Y_train, Y_valid = train_test_split(
    features, target, test_size=0.1, random_state=2022)
print(X_train.shape, X_valid.shape)
```

```
(908, 3) (101, 3)
```

```
models = [LogisticRegression(), SVC(
    kernel='poly', probability=True)]
```

```
for i in range(3):
    models[i].fit(X_train, Y_train)
```

```

    print(f'{models[i]} : ')
    print('Training Accuracy : ', metrics.roc_auc_score(Y_train,
models[i].predict_proba(X_train)[: ,1]))
    print('Validation Accuracy : ', metrics.roc_auc_score(Y_valid,
models[i].predict_proba(X_valid)[: ,1]))
    print()

```

```

LogisticRegression() :
Training Accuracy :  0.540628417292411
Validation Accuracy :  0.5725741780272654

```

```

SVC(kernel='poly', probability=True) :
Training Accuracy :  0.5316444314840465
Validation Accuracy :  0.6433440256615879

```

```

-----
-----
IndexError                                Traceback (most recent call
last)
Input In [39], in <cell line: 4>()
      1 models = [LogisticRegression(), SVC(
      2     kernel='poly', probability=True)]
      4 for i in range(3):
----> 5     models[i].fit(X_train, Y_train)
      7     print(f'{models[i]} : ')
      8     print('Training Accuracy : ', metrics.roc_auc_score(Y_train,
models[i].predict_proba(X_train)[: ,1]))


```

IndexError: list index out of range

```

metrics.plot_confusion_matrix(models[0], X_valid, Y_valid)
plt.show()

```



Precesion of the code is not that good as it's only 37%.